

Rust in Data Science! 🦀
An Outlook of Data Transformation Using Polars

Luciano Scarpulla

Intro

- Hi there! I'm Luciano (lu-cha-no)
- I'm passionate about data and programming
- I love learning new things and sharing knowledge with others
- Currently working for a consultancy company in aviation
- Actively working in open source and exploring Rust to enhance efficiency

Code for this presenetation is available on GitHub



The state of Data science in Rust

"Are we data science yet?"

The screenshot displays the crates.io website for four Rust libraries. Each library entry includes its name, version, download statistics, and a brief description.

- arrow** (crate · repo · docs)
crates.io v54.0.0 downloads 3.1M recent downloads 3.1M Star 2.7k
Rust implementation of Apache Arrow
Last Commit: 2025-01-19 Last Published: 2025-01-18
- ndarray** (crate · repo · docs)
crates.io v0.15.1 downloads 2.0M recent downloads 2.0M Star 3.7k
An n-dimensional array for general elements and for numerics. Lightweight array views and slicing; views support chunking and splitting.
Last Commit: 2024-12-20 Last Published: 2024-08-14
- polars** (crate · repo · docs)
crates.io v0.45.1 downloads 2.3M recent downloads 2.3M Star 81k
DataFrame library based on Apache Arrow
Last Commit: 2025-01-19 Last Published: 2024-12-09
- sprs** (crate · repo · docs)
crates.io v0.11.3 downloads 2.0M recent downloads 2.0M Star 440
A sparse matrix library

The state of Data science in Rust

- Currently there are not many crates for data science in Rust
- Arrow - The universal columnar format and multi-language toolbox for fast data interchange and in-memory analytics
- **Polars** - Dataframes powered by a multithreaded, vectorized query engine, written in Rust

Some exciting features about Polars

Polars is one of the fastest data science tools that is written in Rust with its own compute and buffer implementations, while maintaining compatibility with Apache Arrow. It is more memory efficient than Pandas.

- Native columnar data storage
- Lazy API
- Multi-threaded out of the box
- Cheap Copy-on-Write
- Query plan optimization
 - Predicate pushdown
 - Projection pushdown

Dataframes and Column oriented storage

A Dataframe in polars is a collection of Series (columns) of equal length. Each column has a name and a single data type.

timestamp	id	name	age
---	---	---	---
date	i64	str	i32
2023-01-01	1	John Doe	25
2023-01-02	2	Jane Doe	30
2023-01-03	3	Bob Smith	45

Column oriented storage

Columnar storage is efficient because:

- Better compression - Similar values stored together compress better
- Cache efficiency - Accessing columns keeps data local in memory
- Query optimization - Only reading needed columns reduces I/O
- SIMD operations - Vector operations can process entire columns at once
- Memory efficiency - No need to read unused columns

Column oriented storage

	session_id	timestamp	source_ip
Row 1	1331246660	3/8/2012 2:44PM	99.155.155.225
Row 2	1331246351	3/8/2012 2:38PM	65.87.165.114
Row 3	1331244570	3/8/2012 2:09PM	71.10.106.181
Row 4	1331261196	3/8/2012 6:46PM	76.102.156.138

Traditional Memory Buffer

	1331246660
Row 1	3/8/2012 2:44PM
	99.155.155.225
	1331246351
Row 2	3/8/2012 2:38PM
	65.87.165.114
	1331244570
Row 3	3/8/2012 2:09PM
	71.10.106.181
	1331261196
Row 4	3/8/2012 6:46PM
	76.102.156.138

Arrow Memory Buffer

	1331246660
session_id	1331246351
	1331244570
	1331261196
	3/8/2012 2:44PM
timestamp	3/8/2012 2:38PM
	3/8/2012 2:09PM
	3/8/2012 6:46PM
	99.155.155.225
source_ip	65.87.165.114
	71.10.106.181
	76.102.156.138

```
SELECT * FROM clickstream  
WHERE session_id = 1331246351
```



Lazy vs Eager

- **Eager Evaluation**

```
graph LR; A[ ] --> B[Load Data]; B --> C[Filter]; C --> D[GroupBy]; D --> E[ ]; E --> F[Sort]; F --> G[Result];
```

—► Load Data —► Filter —► GroupBy —►
Sort —► Result

Executed immediately step by step

- **Lazy Evaluation**


```
graph LR; A[Load Data] --- B[ ]; B --- C[Filter] --- D[ ]; D --- E[GroupBy] --- F[ ]; F --- G[Optimize] --> H[Execute] --> I[ ]; G --- J[Result]; H --- J; I --- J; J --- K[Sort];
```

Load Data
Filter
GroupBy
Result
Sort

—► Optimize —► Execute —►

Plan built first, then optimized & executed

Show time!: Buildings in Hong Kong




```

{
  "type": "FeatureCollection",
  "name": "Buildings_in_Hong_Kong",
  "crs": { "type": "name", "properties": { "name": "urn:ogc:def:crs:OGC:1.3:CRS84" } },
  "features": [
    { "type": "Feature",
      "properties": {
        "OBJECTID": 1, "BUILDINGSTRUCTUREID":
5243561, "BUILDINGCSUID": "0162608928T20071224
", "BUILDINGSTRUCTURETYPE": "T", "CATEGORY": "5
", "STATUS": "A", "STATUSDATE": null, "
OFFICIALBUILDINGNAMEEN": null, "
OFFICIALBUILDINGNAMETC": null, "
NUMABOVEGROUNDSTOREYS": null, "
NUMBASEMENTSTOREYS": null, "TOPHEIGHT": null, "
BASEHEIGHT": null, "GROSSFLOORAREA": null, "
RECORDCREATIONDATE": "2007-12-24T00:00:00Z", "
RECORDUPDATEDDATE": "2008-01-16T00:00:00Z", "
SHAPE__Area": 16.947265625, "SHAPE__Length":
18.893572763678002 },
      "geometry": { "type": "Polygon", "coordinates
": ... } },
    ...]

```

Eager API

```
use polars::prelude::*;
use polars_demo::{load_data, unnest_df};

fn main() -> Result<(), Box<dyn
std::error::Error>> {
    let path = std::env::temp_dir().join("
hk_buildings.json");
    load_data(&path)?;
    let file = std::fs::File::open(path)?;
    let mut df = JsonReader::new(file).finish
    ()?.select(["features"])?;
    df = unnest_df(&df)?;
    println!("{:?}", df);
    println!("{:?}", df.column("
RECORDCREATIONDATE")?);
}
```

Eager API

Let's find out how many building records were created over the years!

Eager API

We can create or modify existing columns!

```
df.with_column(  
    df.column("RECORDCREATIONDATE")?  
        .str()?  
        .as_datetime(  
            Some("%FT%H:%M:%SZ"),  
            TimeUnit::Nanoseconds,  
            true,  
            false,  
            None,  
            &ambiguous,  
        )?  
        .year()  
        .with_name("creation_year".into()),  
)?;  
  
println!("{:?}", df.column("creation_year")?);
```


Eager API

```
println!(  
    "{:?}",  
    df.groupby(["creation_year"])?  
        .select(["OBJECTID"])  
        .count()?  
        .sort(  
            ["OBJECTID_count"],  
            SortMultipleOptions::default().  
with_order_descending(true)  
        )  
);
```

Eager API

creation_year	OBJECTID_count
---	---
i32	u32
2005	204182
2008	74603
2007	27891
2009	10094
2006	10088
...	...
2015	6010
2016	4900
2020	4804
2004	1567
2022	1164

How can we make this faster (Lazy API)


The Lazy API in Polars can provide significant performance improvements over eager evaluation by:

- Building a query plan first instead of executing operations immediately
- Allowing the query optimizer to analyze and improve the entire plan
- Pushing predicates down to minimize data read
- Minimizing intermediate allocations
- Parallelizing operations when possible
- Only loading required columns from disk
- Combining multiple operations into optimized steps

How can we make this faster (Lazy API)

```
lf = lf.with_column(
  col("RECORDCREATIONDATE")
    .str()
    .strptime(
      DataType::Datetime(TimeUnit::Milliseconds,
        None),
        StrptimeOptions {
          format: Some("%FT%H:%M:%SZ".
into()),
          strict: false,
          exact: true,
          cache: false,
        },
        lit("raise"),
      )
    .dt()
    .year()
    .alias("creation_year"),
)
.group_by(["creation_year"])
.agg([col("OBJECTID").count()])
.sort(
```

Centroid



We can add native Rust function to our Polars code!

```
use geo::{Centroid, Polygon};

fn calculate_centroid(coords: Vec<Vec<f64>>) ->
Option<(f64, f64)> {
    let polygon = Polygon::new(
        geo::LineString::from(
            coords
                .into_iter()
                .map(|coord| (coord[0], coord[1]
]))),
        .collect:::<Vec<_>>(),
    ),
    vec![],
);

    polygon
        .centroid()
        .map(|centroid| (centroid.x(),
centroid.y()))
}
```

Why Polars in Rust for Production?

- **Performance & Resources**
 - Blazing fast query execution
 - Memory efficient with zero-copy operations
 - Native multi-threading support
 - Perfect for resource-constrained environments
- **Production Ready**
 - Strong type system prevents runtime errors
 - Excellent error handling with Result type
 - Cross-platform compatibility
- **Backend Integration**

```
// Example with Actix-web
use actix_web::{get, web, App, HttpServer,
Result};
use polars::prelude::*;

#[get("/buildings/stats")]
async fn building_stats() ->
Result<web::Json<Value>> {
    let lf = LazyFrame::scan_parquet(
        "data.parquet",
        ScanArgsParquet::default()
    )?
```

Why Polars in Rust for Production?

- **Versatile Data Pipeline Building**
 - Read/Write multiple formats (Parquet, CSV, JSON)
 - Easy integration with Arrow ecosystem
 - Stream processing capabilities
 - Excellent for ETL workflows
- **Developer Experience**
 - Familiar DataFrame API
 - Great documentation and growing community
 - IDE support with type hints
 - Easy to test and benchmark
- **Real-world Use Cases**
 - Data APIs and Microservices
 - Real-time analytics
 - ETL Pipelines
 - Machine Learning Feature Engineering
 - IoT Data Processing

Summary 🚀

- We explored Polars - a powerful DataFrame library for Rust
- Learned about key features:
 - Column-oriented storage
 - Lazy evaluation
 - Native multi-threading
 - Production-ready capabilities
- Demonstrated real-world examples
 - Data loading and transformation
 - Geographic calculations
 - API integration
- The Rust data science ecosystem is growing!

